

Chia-Hsiang Kao

Ph.D. in Computer Science, Cornell University · Expected August 2027

ck696@cornell.edu
iandrover.github.io
Google Scholar · GitHub

RESEARCH INTERESTS

I build world models that infer and predict the hidden dynamics of real-world interactions: recovering 3D motion and physics from observation, and generating controllable video.

- **3D & physical world modeling:** completing occluded scene geometry and motion (amodal scene flow); recovering rigid-body physics from video (Δ ynamics).
- **Video generation & in-context learning:** unifying video tasks through spatiotemporal analogy (ViGeo).
- **LLM agents:** code-executing, tool-using agents and their evaluation (UnivEarth).

EDUCATION

Ph.D. in Computer Science | Cornell University, Ithaca, NY 2023 – Aug 2027 (expected)
Advisor: Bharath Hariharan

Doctor of Medicine (M.D.) | National Yang-Ming Chiao-Tung University, Taipei, Taiwan 2015 – 2022

RESEARCH EXPERIENCE

Student Researcher | Meta, New York, NY Jan 2026 – May 2027
Host: Menglin Jia

- **Amodal scene flow completion (ongoing):** leading scene-level completion of occluded geometry together with its 3D motion, beyond prior visible-only scene flow; designed a novel point-cloud tokenizer and a diffusion model that completes point-voxel tokens while predicting the motion field.
- **ViGeo (Preprint):** led a video in-context learning framework that unifies video tasks as spatiotemporal canvas completion; developed a single model that absorbs a diverse task taxonomy spanning camera control, modality-guided generation, and video manipulation, and generalizes to unseen tasks and sensing modalities; fully fine-tuned Wan2.2-5B on 32 A100 nodes.

Applied Scientist Intern | Amazon, Sunnyvale, CA May 2025 – Aug 2025
Host: Chien-Yi Wang

- **Dynamics (CVPR 2026):** led a VLM that recovers rigid-body physics from monocular video by generating MuJoCo scene configurations as structured text. Trained on 400K synthetic videos with natural-language motion reasoning as auxiliary supervision, it substantially outperforms frontier VLMs, improves further with test-time search, and transfers to a new benchmark of 235 real-world videos.

Graduate Researcher | Cornell University, Ithaca, NY 2023 – Present
Advisor: Bharath Hariharan

- **UnivEarth (ACL 2026 Findings):** led a benchmark of 140 Earth-observation questions across 17 satellite sensors for code-executing LLM agents using the Google Earth Engine API; showed frontier agents reach only 33% accuracy, with generated code failing 58% of the time.
- **AllClear (NeurIPS 2024 D&B):** led a dataset and benchmark for cloud removal in satellite imagery.
- **Counter-current learning (NeurIPS 2024):** proposed a biologically plausible alternative to backpropagation using anti-parallel feedforward and feedback networks.

MENTORING

- Hangyu Zhou: AllClear (NeurIPS 2024 Datasets & Benchmarks, co-first author)
- Cheryl Lam and Aarush Umap: UnivEarth (ACL 2026 Findings)

SKILLS

Modeling: world models, video generation (flow matching, Wan2.2), VLMs, LLM agents & tool use

Tools: PyTorch, FastVideo, multi-node distributed training (32 A100 nodes), MuJoCo, Google Earth Engine

PUBLICATIONS

Google Scholar: 500+ citations, h-index 7 (Oct 2026)

- [S.2] **Unifying Video Tasks via Spatiotemporal Analogy**
Chia-Hsiang Kao, Belinda Zeng, Bharath Hariharan, Menglin Jia
Under submission to ICLR 2027 [arXiv] [Project]
- [C.9] **Dynamics: Language-Based Representation for Inferring Rigid-Body Dynamics From Videos**
Chia-Hsiang Kao, Cong Phuoc Huynh, Chien-Yi Wang, Noranart Vesdapunt, Stefan Stojanov, Bharath Hariharan, Oleksandr Obiednikov, Ning Zhou
CVPR 2026 [arXiv] [Project]
- [C.8] **Towards LLM Agents for Earth Observation**
Chia-Hsiang Kao, Wenting Zhao, Cheryl Lam, Aarush Umap, Shreelekha Revankar, Samuel Speas, Snehal Bhagat, Rajeev Datta, Cheng Perng Phoo, Utkarsh Mall, Carl Vondrick, Kavita Bala, Bharath Hariharan
ACL 2026 Findings; ICML 2025 TerraBytes Workshop [arXiv]
- [C.7] **Counter-Current Learning: A Biologically Plausible Dual Network Approach for Deep Learning**
Chia-Hsiang Kao, Bharath Hariharan
NeurIPS 2024 [arXiv] [Code]
- [C.6] **AllClear: A Comprehensive Dataset and Benchmark for Cloud Removal in Satellite Imagery**
Hangyu Zhou*, Chia-Hsiang Kao*, Cheng Perng Phoo, Utkarsh Mall, Bharath Hariharan, Kavita Bala
NeurIPS Datasets and Benchmarks Track 2024
- [C.5] **Caduceus: Bi-Directional Equivariant Long-Range DNA Sequence Modeling**
Yair Schiff, Chia-Hsiang Kao, Aaron Gokaslan, Tri Dao, Albert Gu, Volodymyr Kuleshov
ICML 2024
- [C.4] **Advancing DNA Language Models: The Genomics Long-Range Benchmark**
Chia-Hsiang Kao*, Evan Trop*, McKinley Polen*, Yair Schiff*, Bernardo P. de Almeida, Aaron Gokaslan, Thomas Pierrot, Volodymyr Kuleshov
AAAI 2023 Workshop; ICLR 2024 MLGenX Workshop
- [S.1] **FedBug: A Bottom-Up Gradual Unfreezing Framework for Federated Learning**
Chia-Hsiang Kao, Yu-Chiang Frank Wang
Preprint
- [C.3] **MEG-Based Classification and Grad-CAM Visualization for Major Depressive and Bipolar Disorders with Semi-CNN**
Chun-Chih Huang, Intan Low, Chia-Hsiang Kao, Chuan-Yu Yu, Tung-Ping Su, Jen-Chuen Hsieh, Yong-Sheng Chen, Li-Fen Chen
IEEE EMBC 2022 [Paper]
- [C.2] **MAML Is a Noisy Contrastive Learner in the Classification**
Chia-Hsiang Kao, Wei-Chen Chiu, Pin-Yu Chen
ICLR 2022
- [C.1] **Demystifying T1-MRI to FDG18-PET Image Translation via Representational Similarity**
Chia-Hsiang Kao, Yong-Sheng Chen, Li-Fen Chen, Wei-Chen Chiu
MICCAI 2021

*: equal contribution

HONORS & AWARDS

- | | |
|--|------------|
| • Foundational AI PhD Fellowship, Cornell University | 2025 |
| • Student Travel Award, MICCAI | 2021 |
| • Undergraduate Research Fellowship, National Science and Technology Council, Taiwan | 2018, 2020 |
| • Summer Research Fellowship, National Health Research Institutes, Taiwan | 2018 |

SERVICE

Conference Reviewer: NeurIPS (2021, 2024–2026), ICLR (2025–2027), ICML (2025), ECCV (2026), COLM (2026), AAAI (2025), AISTATS (2025), AutoML (2022)

Journal Reviewer: CVIU (2022), Computers & Electrical Engineering (2024), IEEE TETCI (2024)